

Supplementary methods

A configurable generative transcriptomics framework for spaceflown mice

Jason Trinh

This document records data roles, model configurations, evaluation gates, statistical safeguards, and supporting results.

S1. OSDR discovery and expression ingestion

OSDR data came from the [Biological Data API](#). Eligible profiles were *Mus musculus* bulk RNA-seq with a resolvable flight or ground-control label, processed RSEM expected counts, and tissue or material metadata sufficient for canonicalization.

The API returned 1,631 rows; aggregation of 21 technical replicates left 1,610 biological profiles (835 flight and 775 ground control) from 75 accessions and 24 canonical material classes.

S2. ARCHS4 cohort audit

The source file, `assets/archs4/mouse_gene_v2.5.h5`, contained 997,515 profiles and 53,511 genes. Three reference cohorts were eligible:

Cohort	Profiles	GEO series	Intended use
Control only	23,614	3,213	Conservative sensitivity cohort
Healthy preferred	62,299	5,307	Primary pretraining source
Broad	134,250	15,111	Diversity sensitivity cohort

Nine GEO series linked to eligible OSDR accessions were excluded before selection. This removed 108 ARCHS4 profiles, including all 53 selected profiles from GSE152382 (OSD-457). The paper-parity run selected 17,244 healthy-preferred profiles across 20 tissues. Complete GEO-series grouping assigned 10,150 profiles to training, 2,466 to validation, and 4,628 to test without series overlap.

S3. Landmark panel and normalization

Full-transcriptome TPM was calculated with GENCODE vM39 mouse gene lengths. Landmark selection occurred after TPM calculation and used a deterministic 974-gene mouse mapping of the human L1000 panel. Training-set MaxAbs scaling was applied after landmark selection.

S4. ARCHS4 DDIM configuration

The full configuration is `configs/rna_diffusion/archs4_mouse_paper_parity_osdr_disjoint.yaml`. Architecture and optimization matched the paper-parity implementation; only reference exclusion and grouped splitting changed.

Component	Value
Input genes	974
Hidden layers	8,192; 8,192
Parameters	227,109,786
Dropout	0.1
Diffusion steps	1,000
Beta schedule	Quadratic, 0.0001 to 0.02
Objective	Summed noise MSE
Optimizer	Adam
Learning rate	0.0004783833151836702
Scheduler	OneCycle
Batch size	2,048
Epochs / optimizer steps	15,000 / 75,000
AMP	Enabled
EMA	0.999
Device	NVIDIA A100-SXM4-40GB
Runtime	6,083 seconds
Peak allocated GPU memory	5.93 GB

S5. Factorized OSDR adaptation

The accepted configuration, `configs/rna_diffusion/osdr_factorized_study_lora512_correlation_refine_osdr_disjoint.yaml`, used the Section S4 backbone. A base adapter received 12,000 domain and 4,000 condition steps before the correlation-refinement stage below. All-tissue and muscle-group screens used the refined adapter.

Component	Value
Backbone	Completed ARCHS4 paper-parity DDIM
Conditioning	Tissue, FLT/GC, accession, material type
Domain adapter	LoRA rank 512, alpha 512
Domain stage	4,000 steps, LR 0.00002
Condition stage	1,000 steps, LR 0.00002
Batch size	512
Condition dropout	0.15
Correlation regularization	Weight 10, 256 genes, timestep ≤ 200
Sampling	100 DDIM steps for evaluation

The split contained 781 training, 536 validation, and 293 locked-test profiles. Every test accession appeared in training, so results measure within-study interpolation. The calibrator used train-only global

alignment, hierarchically shrunk accession and tissue means, positive missing-covariance residual noise, no condition-specific fit of means or covariance, and zero clipping for downstream expression.

S6. Distribution and condition gates

Four generation seeds (5020-5023) were declared before opening the locked test. Metrics were gated independently.

Metric	Gate
Gene-correlation agreement	Paper target ≥ 0.98 ; finite-sample real-bootstrap floor also reported
Precision	≥ 0.95
Recall	≥ 0.85
F1	≥ 0.90
Adversarial accuracy	0.40 to 0.60
FD / real-split P95	≤ 1.0
Pooled effect recovery	Correlation ≥ 0.30 and direction ≥ 0.55
Muscle accession recovery	Correlation ≥ 0.30 and direction ≥ 0.55
Memorization	Generated fraction below train LOO P01 ≤ 0.05

Exact repeats are in `source_data/table_s2_locked_ddim_repeats.tsv`. Locked-test means were correlation 0.9744, precision 0.9974, recall 0.9957, F1 0.9966, adversarial accuracy 0.4753, and FD/real-P95 0.0740. All repeats exceeded the finite-sample correlation floor of 0.9497, although the mean missed the paper target of 0.98. Pooled FLT/GC recovery passed in three of four repeats and muscle accession recovery in all four. The preceding validation screen missed its stricter correlation and muscle gates, so broad biological screens remain developmental. Adversarial accuracy denotes an external nearest-neighbor real-versus-synthetic classifier, not the WGAN critic; 0.5 indicates chance separation.

S7. Configurable benchmark and comparator models

The 463-row experiment planner varied inputs, normalization, transformation, scaling, genes, harmonization, training data, accession and tissue scope, FLT/GC and study conditioning, balancing, seed, and synthetic ratio. Lower-cost screens identified branches for full training.

Liver harmonization benchmark

Nine arms used the same 974-gene DDIM, seed, A100 device, and split: 119 training, 50 validation, and 70 unopened locked-test profiles. No arm passed all fidelity, pooled-condition, and accession-effect gates.

Method	Corr.	Precision	Recall	F1	AA	FD/real P95	Fidelity gate	Accession gate
No harmonization, TPM	0.283	0.200	0.960	0.331	0.850	2.052	Fail	Fail
lIangovan study z-score	0.278	0.160	1.000	0.276	0.770	0.716	Fail	Fail
Mentor two-stage z-score	0.348	0.440	1.000	0.611	0.690	0.977	Fail	Fail
ComBat	0.004	0.040	1.000	0.077	0.810	63.185	Fail	Fail
ComBat-seq	0.067	0.020	1.000	0.039	0.850	156.647	Fail	Fail
MBatch Median Polish	0.009	0.020	1.000	0.039	0.930	205.470	Fail	Fail
MBatch Empirical Bayes	0.001	0.000	1.000	0.000	0.850	60.625	Fail	Fail
MBatch ANOVA	-0.003	0.020	1.000	0.039	0.870	44.716	Fail	Fail
MOBER	0.808	0.260	1.000	0.413	0.770	33.311	Fail	Fail

ComBat, ComBat-seq, and MBatch were transductive; MOBER was inductive. None passed the joint criteria. Tables S13-S14 contain full labels and metrics.

WGAN-GP screen

The WGAN used the Viñas et al. topology: 64-dimensional noise, two 256-unit generator and critic layers, five critic updates, gradient-penalty weight 10, and batch size 32. RMSProp used learning rate 0.0005, alpha 0.9, and epsilon 1e-7 for at most 2,000 epochs. Released-code early stopping began at epoch 1, ran every five epochs, and allowed ten failed checks. Across six seeds, mean correlation, precision, recall, F1, adversarial accuracy, and FD/real-P95 were 0.9759, 0.9764, 0.9938, 0.9850, 0.6362, and 0.1439. Pooled FLT/GC recovery passed 6/6 repeats; accession-aware muscle recovery passed 0/6. These are validation metrics; no separate WGAN test metrics were generated. Table S15 contains exact repeats.

S8. Generated-feature workflow

The funnel tested pooled utility, matched tissue-specific classifiers, consensus panels, and real-data association. Pooled classifiers used real, generated, or combined profiles. The primary tissue-specific comparison used all 974 genes in real-only, generated-only, and equal-weight real-plus-generated classifiers. Section S14 gives its complete design and results.

The secondary consensus screen compared `real_only`, `generated_only`, equal-weight `real_plus_generated`, consensus-ranked `guided_real_only`, and `guided_low_weight` with recentered synthetic profiles at 0.05 total weight.

Eight outer splits evaluated real profiles within accession-condition strata; inner loops selected feature count, regularization, and ranking. Eligibility required nonworse balanced accuracy, AUROC, and average precision versus real-only; ties met this rule without a mean gain. Accessions could occur in both partitions, so this tested within-study development rather than study transfer.

Stable genes required selection frequency ≥ 0.50 and sign agreement ≥ 0.75 . Reinforced genes were stable in real-only and selected synthetic-guided arms with matching directions and a real meta-effect. Synthetic-promoted genes were stable only with eligible guidance and had a real meta-effect. These labels classify selection, not novelty; random-effects and LOO tests followed.

Supplementary Table S9. Complete canonical-tissue utility screen.

Tissue	n (FLT/ GC)	Selected arm	BA real/ selected	AUROC real/ selected	AP real/ selected	Status
Adrenal gland	31 (16/15)	Generated only	0.781 / 0.922	0.906 / 0.984	0.918 / 0.988	Eligible improvement
Bone	30 (15/15)	Guided ranking; real fit	0.656 / 0.734	0.797 / 0.828	0.858 / 0.865	Eligible improvement
Bone marrow	20 (10/10)	Generated only	0.969 / 1.000	1.000 / 1.000	1.000 / 1.000	Eligible improvement
Brain	45 (22/23)	Generated only	0.688 / 0.740	0.712 / 0.802	0.731 / 0.788	Eligible improvement
Brown adipose tissue	20 (10/10)	Generated only	1.000 / 1.000	1.000 / 1.000	1.000 / 1.000	Eligible tie
Cecum	16 (8/8)	Real only	0.875 / 0.875	0.938 / 0.938	0.958 / 0.958	Real-only retained
Cerebellum	44 (23/21)	Generated only	0.529 / 0.602	0.621 / 0.696	0.729 / 0.764	Eligible improvement
Colon	45 (23/22)	Real only	0.656 / 0.656	0.677 / 0.677	0.730 / 0.730	Real-only retained
Eye	18 (9/9)	Generated only	0.594 / 0.781	0.562 / 0.781	0.729 / 0.865	Eligible improvement
Heart	42 (21/21)	Generated only	0.594 / 0.719	0.625 / 0.750	0.642 / 0.778	Eligible improvement
Hippocampus	30 (16/14)	Generated only	0.698 / 0.703	0.740 / 0.781	0.849 / 0.864	Eligible improvement
Kidney	111 (56/55)	Guided ranking; 5% synthetic	0.654 / 0.707	0.680 / 0.771	0.683 / 0.799	Eligible improvement
Liver	197 (101/96)	Real only	0.721 / 0.721	0.764 / 0.764	0.773 / 0.773	Real-only retained
Lung	63 (32/31)	Generated only	0.664 / 0.742	0.680 / 0.830	0.684 / 0.833	Eligible improvement
Mammary gland	20 (8/12)	Generated only	0.771 / 0.885	0.854 / 0.958	0.875 / 0.958	Eligible improvement
Optic nerve	29 (19/10)	Guided ranking; 5% synthetic	0.769 / 0.819	0.863 / 0.950	0.950 / 0.983	Eligible improvement
Retina	63 (37/26)	Guided ranking; 5% synthetic	0.587 / 0.723	0.620 / 0.762	0.756 / 0.846	Eligible improvement
Skeletal muscle, pooled	149 (74/75)	Guided ranking; real fit	0.861 / 0.932	0.924 / 0.961	0.908 / 0.945	Eligible improvement
Skin	125 (66/59)	Real + generated	0.671 / 0.756	0.691 / 0.768	0.747 / 0.808	Eligible improvement
Spleen	86 (44/42)	Real + generated	0.517 / 0.648	0.540 / 0.704	0.589 / 0.749	Eligible improvement
Thymus	94 (51/43)	Guided ranking; 5% synthetic	0.733 / 0.844	0.818 / 0.887	0.862 / 0.918	Eligible improvement
White adipose tissue	20 (10/10)	Generated only	1.000 / 1.000	1.000 / 1.000	1.000 / 1.000	Eligible tie

Supplementary Table S6. Complete anatomical muscle-group utility screen.

Tissue	n (FLT/GC)	Selected arm	BA real/ selected	AUROC real/ selected	AP real/ selected	Status
EDL	24 (12/12)	Real only	0.917 / 0.917	1.000 / 1.000	1.000 / 1.000	Real-only retained
Gastrocnemius	25 (10/15)	Guided ranking; 5% synthetic	0.604 / 0.646	0.646 / 0.750	0.731 / 0.798	Eligible improvement
Quadriceps	35 (18/17)	Real only	0.750 / 0.750	0.880 / 0.880	0.916 / 0.916	Real-only retained
Soleus	41 (22/19)	Real + generated	0.925 / 0.963	0.980 / 0.980	0.980 / 0.986	Eligible improvement
Tibialis anterior	24 (12/12)	Real + generated	1.000 / 1.000	1.000 / 1.000	1.000 / 1.000	Eligible tie

Sample counts are shown as total development profiles followed by flight/ground-control counts. Small cohorts and ceiling-level scores remain exploratory even when a synthetic arm is eligible.

S9. Random-effects reporting and LOO sensitivity

For gene `g` and accession `a`, the real effect was $\text{delta}[g,a] = \text{mean}(\text{FLT},g,a) - \text{mean}(\text{GC},g,a)$. Accession effects were combined by random-effects meta-analysis. Within each tissue, Benjamini-Hochberg adjustment was applied to 974 gene-level P values; BH FDR < 0.05 was the primary rule. Generated profiles never entered this model.

The inventory annotates reinforced, synthetic-promoted, real-only, synthetic-selected without real-direction support, and unselected associations. It also reports study count, direction fraction, τ^2 , and I^2 ; mixed directions qualify but do not negate a significant pooled effect. LOO refitted the model after removing each accession. LOO stability required maximum LOO FDR < 0.05 without sign reversal and was a sensitivity label, not an inclusion rule.

S10. Reactome analysis

The official mouse GMT was generated from `ReactomePathways.txt` and `Ensembl2Reactome_All_Levels.txt`, restricted to *Mus musculus*, `R-MMU-*` pathways, and `ENSMUSG*` genes.

Hypergeometric enrichment used the 974-gene landmark panel as background. Benjamini-Hochberg FDR was applied separately by tissue and selected-gene set. Reactome parent and child terms overlap. Counts of significant rows are therefore not counts of independent biological discoveries.

S11. Skeletal-muscle group analysis

The factorized DDIM and three frozen synthetic development views were reused. No additional neural network was trained for the muscle-group screen.

Group	Profiles	Accessions	FLT	GC
EDL	24	2	12	12
Gastrocnemius	25	3	10	15
Quadriceps	35	4	18	17
Soleus	41	3	22	19
Tibialis anterior	24	2	12	12

Table S6 gives the exact utility results. Soleus reinforced FLT-lower *Bdh1*, *Ech1*, *Bnip3*, and *Decr1* and FLT-higher *Tpm1* across all three accessions. All except *Decr1* passed LOO sensitivity; *Mef2c* and *Pxmp2* were not promoted. Gastrocnemius promoted *Fhl2* and *Nfkb1a*; tibialis anterior reinforced *Cdkn1a*, *St3gal5*, and *Bnip3* and promoted *Cebpd*. None passed LOO FDR. EDL and quadriceps retained real-only arms.

LOO refitted only the real-data meta-analysis; it did not remove an accession from generator adaptation. The muscle findings therefore remain developmental pending a fully excluded study.

S12. Random-effects BH-FDR gene inventories

Table S11 contains all real-data random-effects associations with BH FDR < 0.05; Table S12 gives counts for every analysis unit. The 459 tissue-gene rows comprise 202 associations across 10 of 22 canonical tissues and 257 across five muscle groups. They are not independent discoveries because genes and samples can recur across analyses. Direction was unanimous for 363 rows and mixed for 96; this was an annotation, not a filter.

Analysis unit	BH-FDR genes	FLT higher	FLT lower	Unanimous direction	Synthetic-promoted	Reinforced
Adrenal gland	22	1	21	21	1	1
Eye	8	4	4	8	0	1
Heart	4	2	2	4	0	0
Kidney	3	3	0	1	1	1
Liver	19	8	11	0	0	0
Retina	5	5	0	5	0	0
Skeletal muscle, pooled	26	12	14	5	4	8
Skin	3	2	1	1	1	0
Spleen	34	32	2	21	3	1
Thymus	78	37	41	57	13	3
EDL	136	14	122	130	0	0
Gastrocnemius	15	8	7	13	2	0
Quadriceps	29	25	4	22	0	0
Soleus	29	5	24	27	0	5
Tibialis anterior	48	42	6	48	1	3

Table S10 is the 49-row secondary consensus subset: 26 promoted and 23 reinforced results with supporting real effects. Attribution was suppressed when the generated arm failed eligibility. The remaining inventory contains 34 real-only selections, 370 BH-significant unselected results, and six synthetic selections without real-direction support. Below, direction is the real random-effects estimate; an asterisk marks disagreement among accessions.

Synthetic-promoted genes

Analysis unit	FLT higher	FLT lower
Adrenal gland	None	Psmb8
Kidney	Inpp4b*	None
Skeletal muscle, pooled	None	Klhl21*, Mapkapk5*, Reep5*, Itgb5*
Skin	Plscr1	None
Spleen	Rai14, Myl9*, Ptpkr	None
Thymus	Hsd17b11*, Etv1	Nusap1, Stmn1, Birc5, Cdk1, Top2a, Ccnb2, Aurka, Ccne2, Kif20a, Pcna, Ccnf
Gastrocnemius	Nfkbia	Fhl2
Tibialis anterior	Cebpd	None

Reinforced genes

Analysis unit	FLT higher	FLT lower
Adrenal gland	None	Tspan4
Eye	None	Klhl21
Kidney	Slc37a4	None
Skeletal muscle, pooled	Sox4, Cebpd*, Sh3bp5, Prkcd, Arid5b*, Sesn1*, Tle1*	Bph1*
Spleen	Loxl1	None
Thymus	Snx7	Ube2c, Gmn
Soleus	Tpm1	Bdh1, Ech1, Bnip3, Decr1
Tibialis anterior	Cdkn1a, St3gal5, Bnip3	None

Heart, liver, retina, EDL, and quadriceps had BH-FDR genes but none was synthetic-informed. Bone, bone marrow, brain, brown adipose tissue, cecum, cerebellum, colon, hippocampus, lung, mammary gland, optic nerve, and white adipose tissue had no BH-FDR gene in the landmark panel. Tables S11-S12 retain the complete results, including pooled muscle, liver's 19 real-only associations, skin's promoted *Plscr1*, and the null lung inventory.

Machine-readable Tables S1-S21 are under [source_data/](#).

S13. Targeted literature annotation of synthetic-informed genes

The literature screen covered all 49 synthetic-informed tissue-gene associations in Table S10: 26 promoted and 23 reinforced. Searches used the gene symbol and synonyms together with the tissue, spaceflight or microgravity, and relevant process terms. Primary studies were preferred for directional and mechanistic claims. Public OSDR or GeneLab reanalyses were retained as context but were not counted as independent confirmation.

Selection status and literature interpretation were recorded as independent dimensions. `selection_status` is `promoted` when a gene was stable only under synthetic-informed selection and `reinforced` when it was stable under both real-only and synthetic-informed selection. `literature_classification` is one of four mutually exclusive labels. Aligning evidence agrees with the observed gene direction or the same-tissue process. Complementary evidence connects the gene to a

reported flight process or relevant mechanism without reproducing the same gene-tissue-direction result. Ambiguous evidence is mixed or context dependent. Unmatched means that the targeted search found no sufficiently specific tissue or process match. Unmatched is not proof of novelty and does not imply that the gene lacks a plausible mechanism.

Selection status	Aligning	Complementary	Ambiguous	Unmatched	Total
Promoted	11	13	1	1	26
Reinforced	11	6	3	3	23
Total	22	19	4	4	49

Five records used a direct same-gene, same-tissue, same-direction transcript-level scope: thymus `Ccne2`, thymus `Ccne2`, gastrocnemius `Nfkb1a`, and pooled-muscle `Sox4` and `Arid5b`. Soleus `Ech1` and `Decr1` matched the gene/protein, tissue, and direction in flight proteomics and were recorded separately as cross-assay alignment. Process-level alignment covered the remaining flight-lower thymus mitotic genes, reinforced thymus `Ube2c` and `Gmnn`, flight-lower eye `Klhl21`, several pooled-muscle genes, and soleus `Bdh1`.

The four ambiguous records were thymus `Birc5`, soleus `Bnip3`, soleus `Tpm1`, and tibialis-anterior `Cdkn1a`. Their prior directions varied by mission, assay, species, sex, or muscle context. The four unmatched records were adrenal `Psmb8` and `Tspan4`, pooled-muscle `Bph1`, and thymus `Snx7`. These candidates remain statistically associated in the real-data analysis but should not be presented as literature-supported mechanisms.

Table S16 records selection status, literature class, interpretive role, evidence scope, source relationship, and a gene-level interpretation. Table S17 supplies the 33-source bibliography and states whether each source is independent, potentially overlapping, or mechanistic context only. The deterministic builder is `nasa_mouse_rna_diffusion.annotate_promoted_gene_literature`.

S14. Matched all-gene classifier analysis

The primary classifier analysis held the feature space and model choices fixed across training sources. It covered 22 canonical tissues and five skeletal-muscle groups. Each ridge logistic classifier used all 974 landmark genes. Eight outer splits were run per analysis unit, yielding 648 fitted classifiers across real-only, generated-only, and real-plus-generated arms.

The scaler was fitted on real outer-training profiles. Ridge regularization was selected from inner real-only data and then reused for all three arms in that split. The combined arm assigned the full synthetic set the same total weight as the real set. Every arm was scored on identical held-out real profiles. Metrics were calculated both after pooling those profiles and after calculating accession-level scores and averaging them. A synthetic arm passed the joint utility gate when mean balanced accuracy, AUROC, and average precision were all no worse than real only in both summaries.

Accession-blocked permutation importance used ten shuffles per gene and fit. A gene had repeat-consistent marginal importance when its mean held-out-real AUROC loss was at least 0.001 and the loss was positive in at least half of the outer splits. Exact linear SHAP supplied contribution direction relative to the shared real-training background. A biological candidate also required real-data BH FDR below 0.05, a synthetic arm that passed the joint utility gate, a coefficient direction matching the real FLT/GC effect, and positive FLT/GC SHAP separation.

Real-plus-synthetic training passed all six mean metric checks in 18 of 27 units. Sixteen improved at least one metric; bone marrow and brown adipose tissue tied. Synthetic-only training passed in six units, five with at least one improvement and one tie.

Twenty-one unique tissue-gene associations passed the complete matched gate. Two were supported by both synthetic arms, producing 23 arm-level rows.

Tissue	FLT higher	FLT lower	Unique associations
Thymus	Klhdc2, Snx7, Etv1, Plscr1, Tspan3, Socs2	Nusap1, Stmn1, Birc5, Ccnb2, E2f2, Ube2c, Cdc20, Gmnn, Kif20a	15
Liver	None	Grb10, Ppic, H2-DMa, Gtf2a2	4
Skin	Plscr1	None	1
Spleen	Loxl1	None	1

Seven thymus genes were promoted under the matched definition: Klhdc2, Snx7, Etv1, Plscr1, Tspan3, Socs2, and Kif20a. Their marginal importance passed only after real-plus-synthetic training. The remaining candidates retained measurable importance in both real-only and synthetic classifiers. No anatomical muscle group passed the full matched utility and gene-importance gate.

Reactome enrichment of all 15 thymus genes tested 36 terms; 26 had FDR below 0.05. The leading term was mitotic cell cycle (R-MMU-69278, FDR 0.004744), with Ube2c, Kif20a, Cdc20, Gmnn, E2f2, and Ccnb2. Thirty-three terms were significant in the flight-lower subset. The four liver genes tested two terms, neither significant.

The matched and consensus analyses retained 21 and 49 associations, respectively. Eleven were shared, 38 appeared only in the consensus panel analysis, and ten appeared only in the matched analysis. The methods estimate different quantities. The matched workflow measures one gene's marginal importance while all 974 genes remain available. With correlated genes, shuffling one member leaves related genes able to carry the same signal, so the individual permutation loss can be small. Ridge can also divide coefficient weight among those genes. Consensus ranking is less strict at the individual-gene level and can retain several members of one correlated program. Matched results therefore support claims about synthetic contribution; consensus results support secondary panel and pathway interpretation.

Tables S18-S21 contain matched utility, retained candidates, candidate Reactome enrichment, and the matched-consensus crosswalk. The outer profiles were held out from classifier fitting and regularization selection, but the fixed DDIM had seen the broader development pool. This is not an independent held-out-study test. Generated profiles remained model draws and did not enter BH-FDR estimation.

S15. Supplementary figures




Figure S1. Repeated nested muscle-group balanced accuracy. Each row is a muscle group and each column is a downstream use of real or generated profiles. Arm selection also required nonworse AUROC and average precision.

Pooled augmentation failed, while tissue-specific synthetic use varied




Figure S2. Downstream utility of generated expression. (A) Direct pooled augmentation on the locked real-profile test. (B) Selected-arm changes in balanced accuracy, AUROC, and average precision across repeated tissue-specific development splits. All evaluations used real profiles.